

Seaborn Exercises

Instructions:

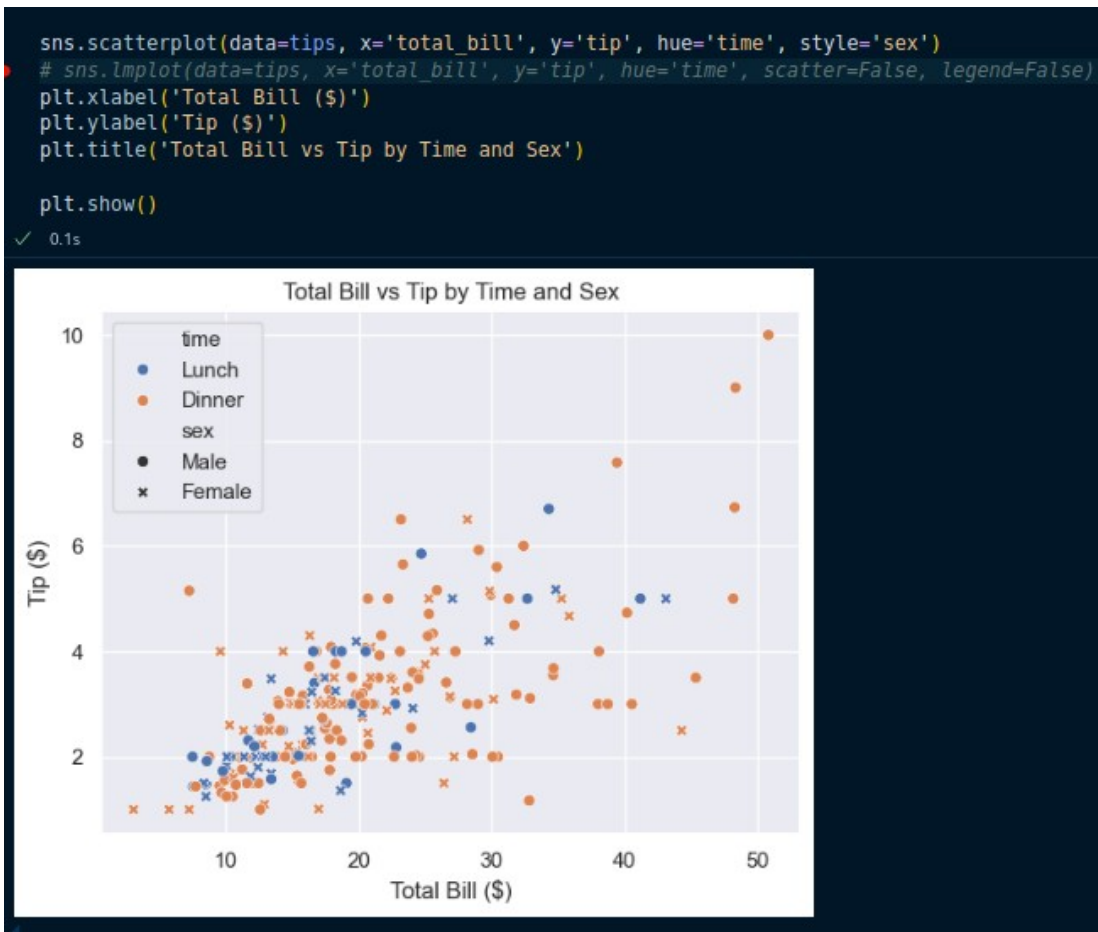
- Complete the tasks using Python (pandas, numpy, seaborn).
- Include code cells, resulting plots, and short interpretations where requested.
- Submit a single notebook with all answers and save final figures as PNG files.
- This worksheet contains 10 exercises arranged from easy → hard.

1. 1) Load & visualize: `tips` scatter

Dataset: `sns.load_dataset('tips')`

Task: Create a scatter plot of `total_bill` vs `tip`. Color points by `time` and use different markers for `sex`. Add axis labels and a title.

Bonus: Add a regression line for each time (lunch/dinner) on the same axes.

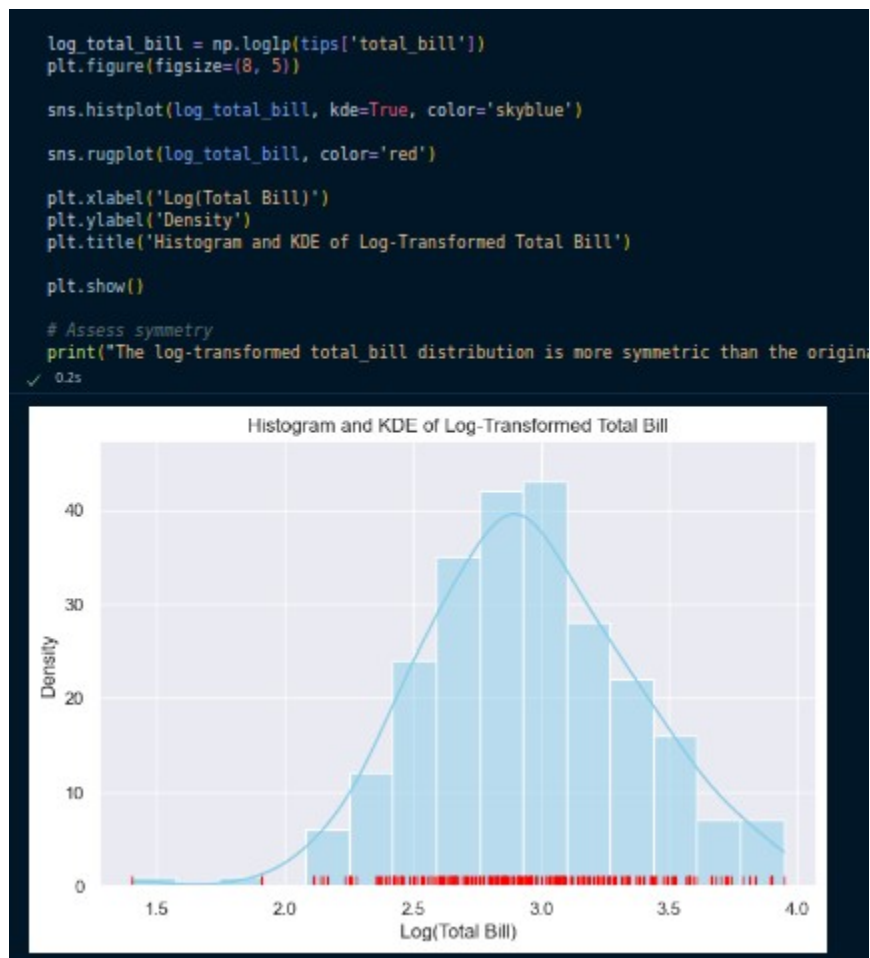


2. 2) Distribution + NumPy: histogram & KDE

Dataset: tips (use total_bill)

Task: Using numpy, compute a log-transformed version of total_bill (np.log1p). Plot its histogram and KDE on the same axes. Briefly state whether the transform made the distribution more symmetric.

Bonus: Overlay a rug plot.



3. 3) Categorical aggregation: barplot with pandas groupby

Dataset: titanic

Task: Using pandas, group by class and compute the survival rate; show a bar plot of survival rate per class. Annotate bars with percentages.

Bonus: Split bars by sex using hue.



4. 4) Pairwise relationships: pairplot vs PairGrid

Dataset: iris

Task:

- Produce `sns.pairplot(iris, hue='species')`.
- Create the same using `PairGrid` with scatterplots in the upper triangle and KDEs on the diagonal. Explain one advantage of `PairGrid`.



5. 5) Time-series & heatmap: flights pivot

Dataset: flights (year, month, passengers)

Task: Pivot into a month x year matrix and plot a heatmap with annotations and colorbar. Interpret the main pattern you observe.

Bonus: Ensure months appear in chronological order.

```

flights = sns.load_dataset('flights')

flights_pivot = flights.pivot(index='month', columns='year', values='passengers')


# print(flights_pivot)
month_order = ['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec']
flights_pivot = flights_pivot.reindex(month_order)

plt.figure(figsize=(10, 6))
sns.heatmap(flights_pivot, annot=True, fmt='d', cmap='YlOrRd', cbar=True)

plt.xlabel('Year')
plt.ylabel('Month')
plt.title('Passenger Numbers by Month and Year')

plt.show()

```



6. 6) Regression with groups & matplotlib tweak

Dataset: mpg (model_year, mpg, origin)

Task: For each origin, plot a linear regression (mpg vs model_year) on the same axes with different line styles. Add a legend outside the plot and a vertical dashed line at the median model year.

Bonus: Compute and display the slope for each origin as text on the plot.

```

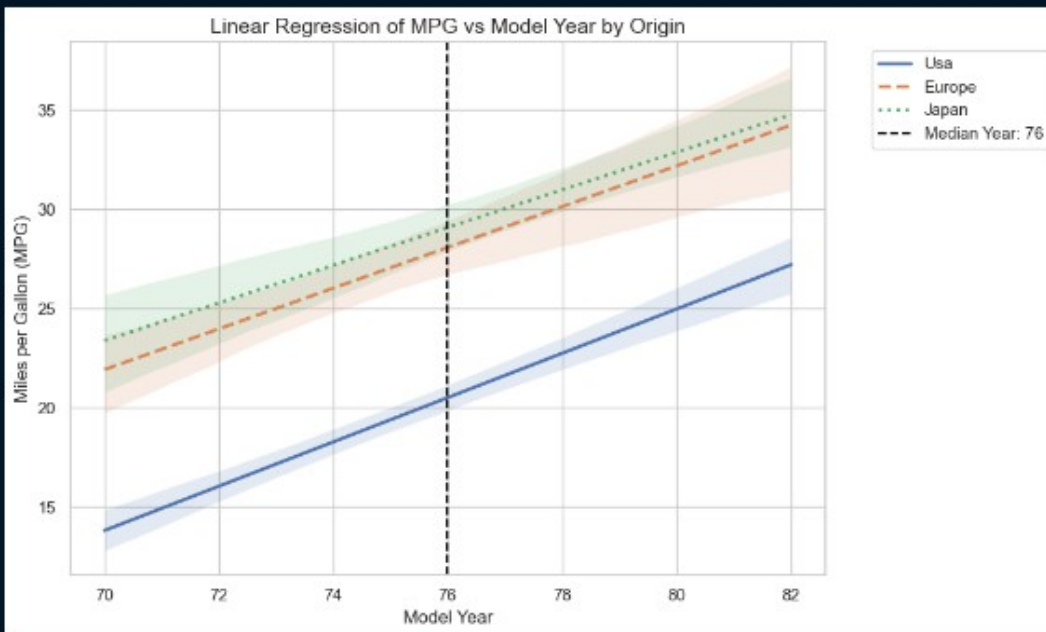
mpg = sns.load_dataset("mpg").dropna(subset=["mpg", "model_year", "origin"])
median_year = int(mpg["model_year"].median())
sns.set(style="whitegrid")
plt.figure(figsize=(10,6))

line_styles = {"usa": "-", "europe": "-.", "japan": ":"}
for origin, style in line_styles.items():
    sns.regplot(
        data=mpg[mpg["origin"] == origin],
        x="model_year",
        y="mpg",
        scatter=False,
        label=origin.capitalize(),
        line_kws={"linestyle": style}
    )

plt.axvline(median_year, color="black", linestyle="--", label=f"Median Year: {median_year}")
plt.xlabel("Model Year", fontsize=12)
plt.ylabel("Miles per Gallon (MPG)", fontsize=12)
plt.title("Linear Regression of MPG vs Model Year by Origin", fontsize=14)

plt.legend(bbox_to_anchor=(1.05, 1), loc="upper left")
plt.tight_layout()
plt.show()

```



7. 7) Faceting & custom aggregation with FacetGrid

Dataset: tips

Task: Create a FacetGrid faceted by day (columns) and smoker (rows). In each facet show a violin plot of total_bill and overlay the mean point (white dot with black edge). Annotate each facet with the mean value.

8. 8) Joint distributions + conditional coloring

Dataset: penguins (or substitute)

Task: Create a jointplot of bill_length_mm vs bill_depth_mm with hex bins. Color hexes by the average species-encoded-as-integer inside each bin (compute binned-statistic). Explain your binning approach.

Bonus: Provide a legend mapping species encoding back to names.

9. 9) Build a custom seaborn helper function

Task: Implement pretty_violin(df, x, y, hue=None, title=None) that:

- sets a publication-style seaborn theme
- draws a violin plot with inner quartiles
- overlays swarm points (max 200 per group, sample if needed)
- adds a text box with mean, median, std for each group

Demonstrate on a synthetic dataset (3 groups).

Bonus: Return fig, ax.

10. 10) Exploratory mini-project

Dataset: choose one (titanic, penguins, mpg, or tips)

Deliverables: For the chosen dataset produce cells that:

1. Clean missing data and describe strategy.
2. Show a correlation heatmap and explain three strongest associations.
3. Compare a numerical outcome across two categorical variables.
4. Create a multi-panel facet visualization that reveals an interaction effect.
5. Create one advanced visualization (annotated regression + residuals OR hexbin + marginals OR clustermap).
6. Summarize your 3 most important findings in bullet points.

Bonus: Add an interactive widget to choose hue/col (optional).